

One page, ready before the laptop opens

A morning triage digest for a founder drowning in 500 emails:
an **agentic CLI** over a **deterministic toolbox** — the model decides what to chase and how to say it; every fact arrives through a cited, pure function it cannot forge.

Go • aubade (product) + aubade-lab (harness, never ships) • zero-key --no-llm path • vendor-neutral Runner (claude CLI now)

Architecture — the split is the opinion

data/ → DataSource (LocalFS now • Composio wk-2) → Event model → 7 cited extractors (commitments • quiet-threads • conflicts • contradictions • dispatchables • suppressions • staleness) → signals.json → agentic compose (citation-validated) | --no-llm template → digest.md → aubade-lab eval

- **Facts are deterministic, cited, byte-stable.** Same corpus + same --today ⇒ identical signals — which is what makes a trap eval possible at all.
- **Orchestration without fabrication.** The runner picks tools and composes the page, but a validator rejects any page citing a ref absent from that run's signals.json .
- **Consensus on by default.** Two bounded decisions (ambiguous urgency; the one-thing pick) are voted by every runner that answers a liveness probe; disagreement lands in "I'm not sure," never a coin flip. A CEO's wrong priority costs more than 3× spend.
- **The honesty floor is not customizable.** --customize reshapes the compose stage only; staleness, contradictions and "I'm not sure" are appended from signals afterward.

Considered and rejected

ALTERNATIVE	WHY IT LOST — AND WHAT WAS KEPT
One big LLM prompt	Unfalsifiable: no signals to assert against; hallucinated confidence indistinguishable from judgment. Rejected on evaluability — the graded axis.
Deterministic-only	Provable but judgment-free; can't decide a 14-message thread needs reading. Kept as --no-llm : the zero-key path and the control the agentic mode is diffed against.
RAG / vector store	500 emails fit in memory; retrieval adds a recall failure mode to a scale problem that doesn't exist.
Omnigent meta-harness	Alpha, Python, session-scale. Kept its one idea — several harnesses, one question — as ~200 lines of consensus over our own Runner.
SDK loop, this week	Costs key handling, buys nothing yet. Runner (Ask + Orchestrate) is vendor-neutral: an SDK runner is a registration, not a rewrite.
Web UI / Slack bot	Explicitly unwanted. CLI in, markdown out; the engine is the product.

Evals — the exam, the student, and the answer key

Each planted trap is a *scenario script* emitting both its data and its traps.json entry — the answer key is a return value of the exam it grades, so they cannot drift. Code graders first; the model judges only what code cannot. 14 traps must surface; 6 must not (suppression is proven, not assumed); a committed golden digest proves every trap was catchable.

20 / 20

regression traps caught (--no-llm), 100% bar, gated in make check

pass³ 20/20

agentic suite: 3 isolated trials, pass³ and pass@3 — never one number

52 / 0

citations on the page / citations failing to ground in signals.json

SABOTAGE: EXTRACTOR OFF	DROP
quiet-threads	-4
conflicts	-3
dispatchables	-2
commitments • contradictions • suppressions • staleness	-1 each

Each extractor disabled in turn: the score must drop, or a non-zero alarm says the graders are blind.

Second week

- **Real data:** Composio provider behind the existing DataSource: only the loader changes.
- **Runners:** SDK-backed runners beside the CLI ones; aubade mcp serve exposes the toolbox to every MCP-aware framework.
- **Adversarial eval:** the LLM authors traps we didn't write, injected into a copy of the corpus — the tool passes exams it never studied for.
- **Anchored judge** (reason-before-score, "uncertain" allowed) grades the one axis code can't: "reads like the sample, in Avery's voice." It informs, never blocks — and it called the template page reads-like-a-machine : the anchors discriminate.
- **Eval depth:** dataset mutation testing — perturb dates and senders; traps must still be caught.
- **Delivery:** GitHub Actions cron 05:45 PT (DST-aware), date-stamped idempotent output, per-step secrets, --no-llm as the degraded mode: an expired token costs a plainer page, never a missing one.